Team Members

Ali Taha

Ahmad Shata

Andrew Anter

Merit Mekhail

Hadi Lotfy

Developer

DevOps Engineer

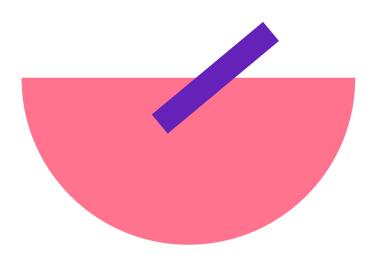
Solution Architect

Scrum Master

Product Owner

ITI - Devops Intensive Training - GCP Course

Airflow DAG in Cloud Composer





Agenda

- Product Owner
 - Product Backlog
- Scrum Master
 - Sprint backlog
- Solution Architect
 - What is composer?
 - Why composer?
- Devops Engineer
 - GCP
 - Terraform
- Developer
 - DAG

Product Backlog

PRODUCT BACKLOG

User Story ID	User Story	Estimate (size)	Priority	Sprint	Estimated effort
US001	As Admin I need to take backup our database periodically So that we can return to same state in case of failure	large	1	1	24h
US002	As Admin I need to delete old backups So that we save storage	small	3	3	7h
US003	As Admin I need to see a list of all members and visitors So that I monitor site visits	large	2	2	24h
US004	As Customer I need to be able to login So that I can post new entries	medium	2	2	15h

Sprint Backlog

SPRINT 1 BACKLOG

ID	User Story	Tasks	Owner	Status	Day 1	Day 2	Day 3	Day 4	Day 5	Day 6	Day 7	Estimated effort
US001	As Admin I need to take backup our database So that we can return to same state in case of failure	Create VPC and Subnet using Terraform	Shata	Done	4							
		Create Composer	Shata	Done		3	4					
		Composer Documentation	Shata	Done				3				24
		Create a DAG on the composer to Automate taking machine image Task	Ali	Done				4	3			
		DAG Documentation	Ali	Done						3		
Total				4	3	4	7	3			24	

Sprint Duration: One Week

What is Composer?



 A platform created by the community to programmatically author, schedule and monitor workflows.



Google Cloud Composer

 A fully managed workflow orchestration service built on Apache Airflow.

Why Composer? Benefits Of Using Cloud Composer

- Multi-cloud
- Open source
- Hybrid
- Integrated
- Python programming language
- Reliability
- Fully managed
- Networking and security

Spinning up our Infrastructure

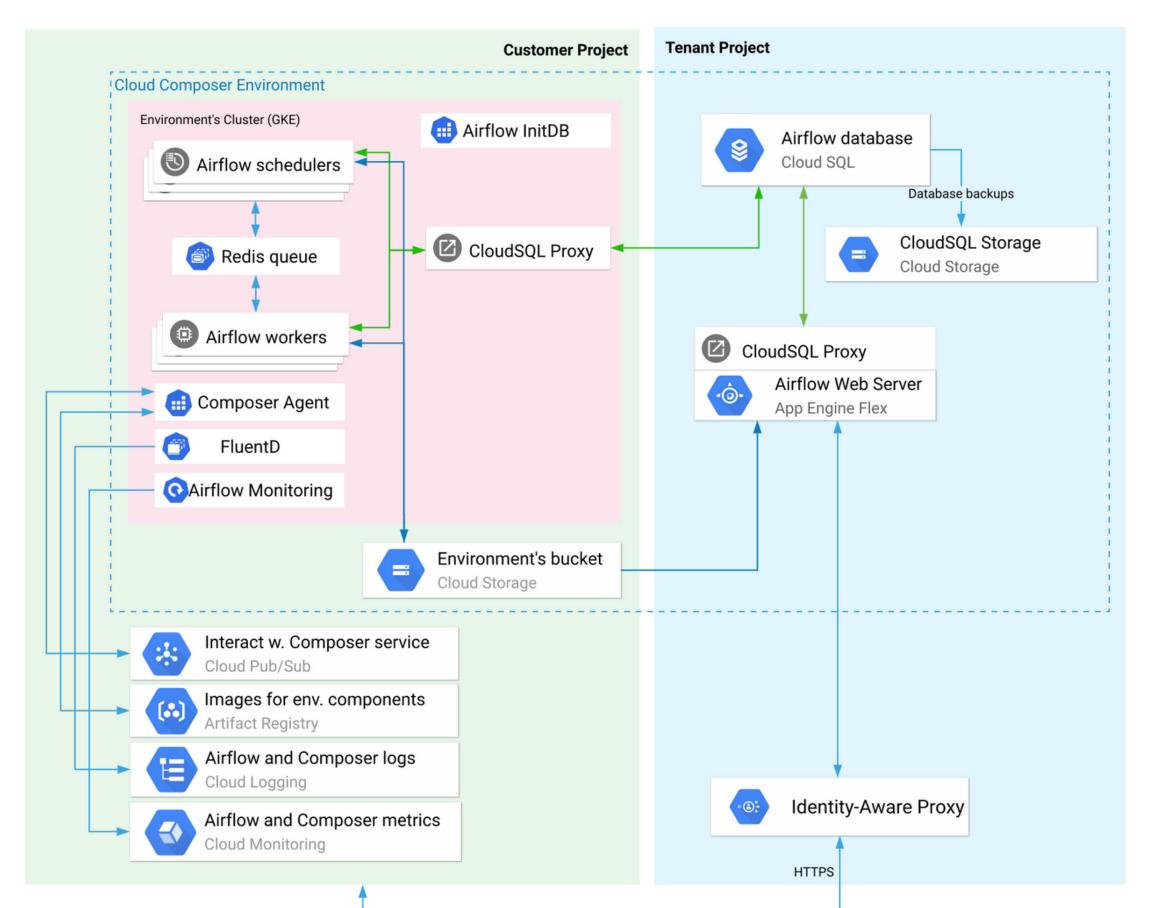
We are going to use Terraform as our IaC utility rather than creating the infrastructure directly from the console to exploit the Terraform idempotency, modulation, and backend



Network Module

```
resource "google_compute_network" "net" {
                                 = "shata-387907"
 project
                                 = var.vpc-name
  name
 auto_create_subnetworks
                                 = false
resource "google_compute_subnetwork" "subnet" {
                                 = var.subnet-name
 name
                                 = "shata-387907"
 project
 private_ip_google_access
                                 = var.private-google-access
                                 = google_compute_network.net.id
 network
                                 = var.subnet-cidr
 ip_cidr_range
 region
                                 = var.subnet-region
  secondary_ip_range {
                                 = var.secondary-ip-range-1-name
   range name
   ip_cidr_range
                                 = var.secondary-ip-range-1-cidr
  secondary_ip_range {
                                 = var.secondary-ip-range-2-name
   range_name
   ip_cidr_range
                                 = var.secondary-ip-range-2-cidr
```

Composer Structure Depiction



Network Module Output

```
output "vpc-uri" {
  value = google_compute_network.net.self_link
}
output "subnet-uri" {
  value = google_compute_subnetwork.subnet.self_link
}
```

Enabling Composer API

Composer Module

```
# The composer env
resource "google_composer_environment" "test" {
 project = var.project-id
 name = var.composer-name
region = var.composer-region
 config {
   environment_size = var.environment-size
   node config {
     network = var.vpc-uri
     subnetwork = var.subnet-uri
     service_account = "${google_service_account.test.email}"
   software_config {
  image_version = var.composer-version
 depends_on = [
   google_project_iam_member.composer-worker,
   google_project_iam_member.composer-ext,
   google_service_account_iam_member.API_service_account
```

Service Account

Least SA privilege

```
# Add a composer worker role to the service account
resource "google_project_iam_member" "composer-worker" {
                         = var.project-id
  project
  role
                          = "roles/composer.worker"
                          = "serviceAccount:${google service account.test.email}"
 member
# Add a Cloud Composer v2 API Service Agent Extensio role to the service account (only required for composer V2)
resource "google_project_iam_member" "composer-ext" {
  project
                         = var.project-id
                          = "roles/composer.ServiceAgentV2Ext"
  role
                          = "serviceAccount:${google service account.test.email}"
  member
```

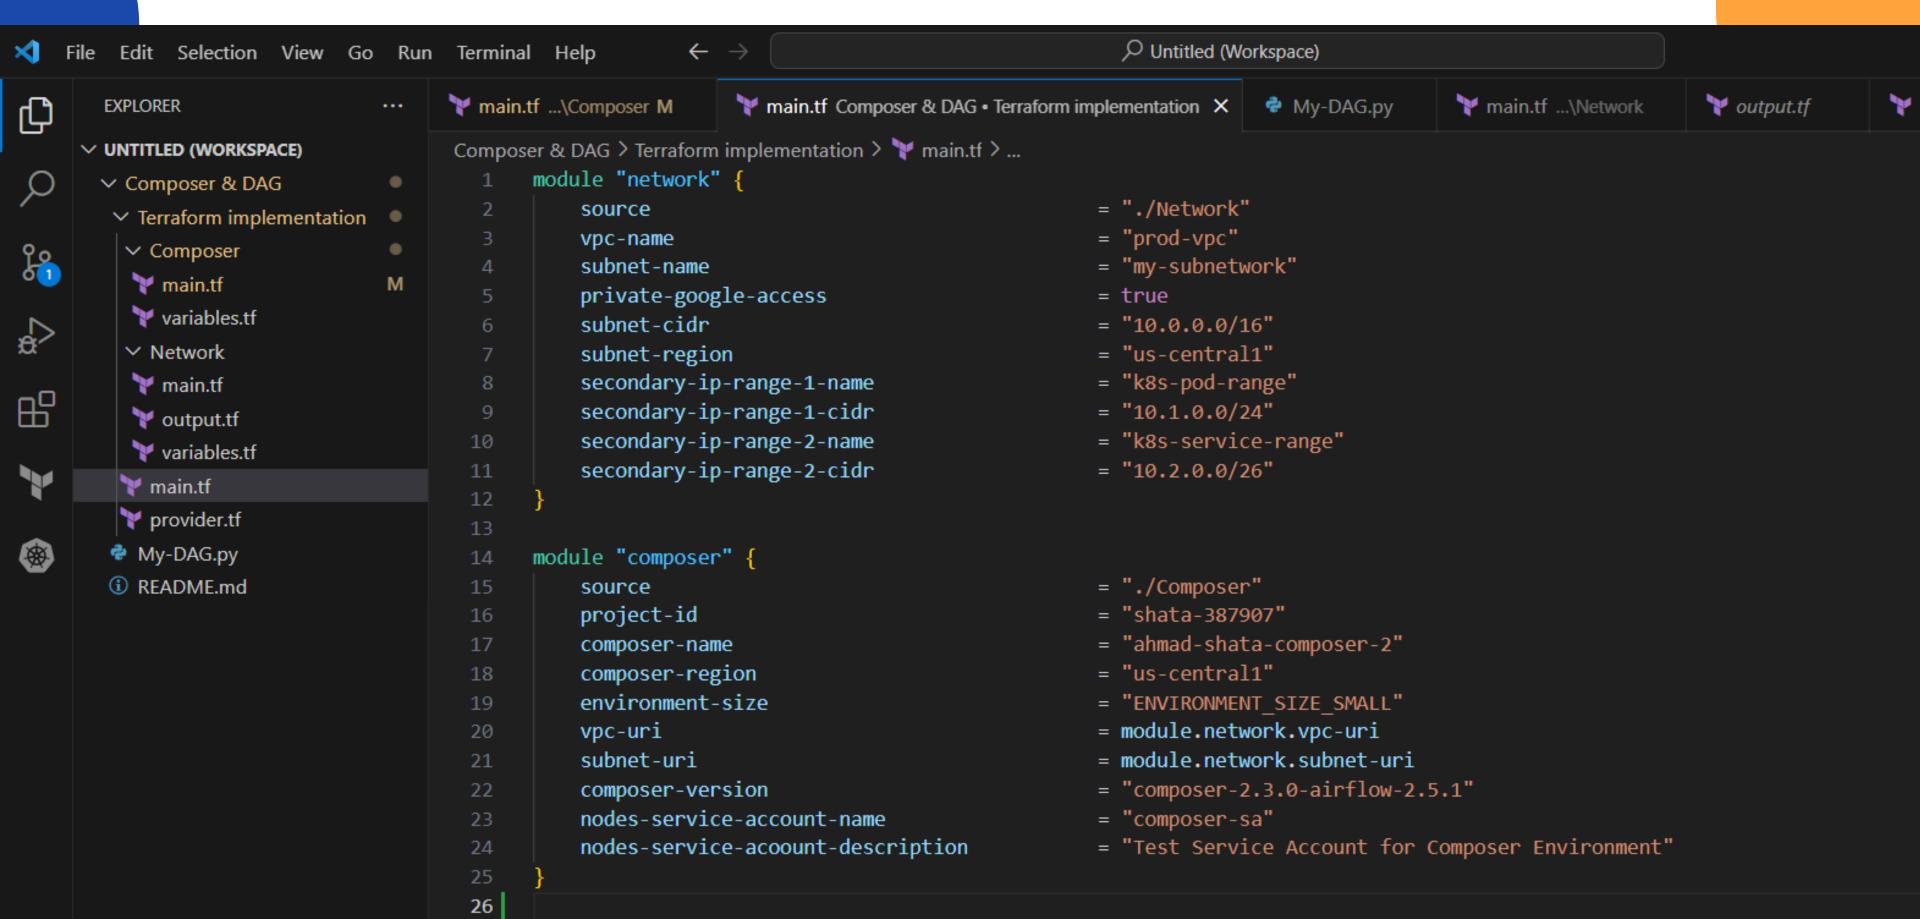
Custom IAM Role

```
resource "google project iam custom role" "my-composer-sa-role" {
 role id = "composer-sa-role"
 title = "My Instance Role"
 description = "my custom iam role"
  permissions = [
    "compute.machineImages.create",
   "compute.disks.createSnapshot",
   "compute.instances.useReadOnly",
    "compute.instances.start",
   "compute.instances.stop"
```

Attaching the custom role

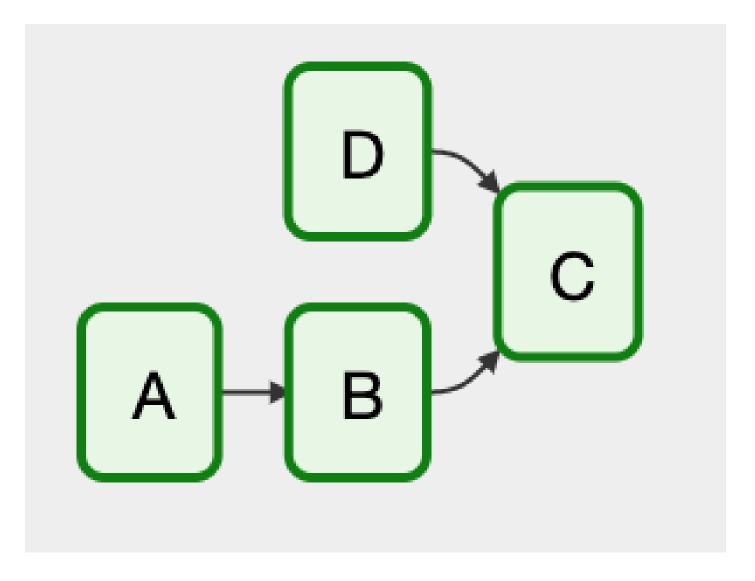
Adjusting Service Agent

Calling our Modules



WHAT IS DAG

A DAG (Directed Acyclic Graph) is the core concept of Airflow, collecting Tasks together, organized with dependencies and relationships to say how they should run





```
default_args = {
    'start_date': airflow.utils.dates.days_ago(0),
    'retries': 0
}
```

- The default_args dictionary contains default configuration parameters for the DAG.
- The 'start_date': airflow.utils.dates.days_ago(0) specifies the start date for the DAG as the current date.
- The 'retries': 0 specifies that the DAG should not be retried in case of failure. Change this value if you want to allow retries.

```
dag = DAG(
   'backup',
   default_args=default_args,
   description='backup',
   schedule_interval='0 3 * * 5',
   max_active_runs=2,
   catchup=False,
   dagrun_timeout=timedelta(minutes=10),
)
```

- · The dag variable creates an instance of the DAG class.
- The 'backup' parameter is the DAG's identifier.
- The default_args=default_args sets the default configuration parameters for the DAG.
- · The 'backup' string in the description parameter provides a brief description of the DAG.
- The '0 3 * * 5' schedule_interval parameter sets the schedule to run the DAG every Friday at 3 AM.
- The max_active_runs=2 parameter specifies that a maximum of two DAG runs can be active at the same time.
- The catchup=False parameter prevents Airflow from running any missed or overdue DAG runs when the DAG is first created.
- The dagrun_timeout=timedelta(minutes=10) parameter sets a timeout duration of 10 minutes for each DAG run.

```
stop_vm = BashOperator(
    task_id='stop_vm',
    bash_command='gcloud compute instances stop instance --zone=us-central1-a',
    dag=dag,
    depends_on_past=False,
    priority_weight=2**31 - 1,
    do_xcom_push=False
)
```

- The stop_vm variable creates an instance of the BashOperator class to stop a virtual machine instance.
- The 'stop_vm' parameter sets the task's identifier.
- The 'gcloud compute instances stop instance --zone=us-central1-a' bash_command parameter specifies the Bash command to execute. It stops a
 virtual machine instance named 'instance' in the 'us-central1-a' zone using the gcloud command.
- The dag=dag parameter assigns the task to the defined DAG.
- The depends_on_past=False parameter indicates that the task does not depend on the past successful runs of other tasks.
- The priority_weight=2**31 1 parameter sets the priority weight of the task to the maximum value.
- The do_xcom_push=False parameter disables the pushing of task output to XCom.

stop_vm >> create_machine_image >> start_vm

The >> operator connects the tasks to define their dependencies. Here, the stop_vm task must complete successfully before the
create_machine_image task can start, and the create_machine_image task must complete successfully before the start_vm task can start.

